PREPARED FOR

Emirates

SUBMITTED BY

Aonami

EMIRATES GROUP · HR COMPLIANCE CHALLENGE

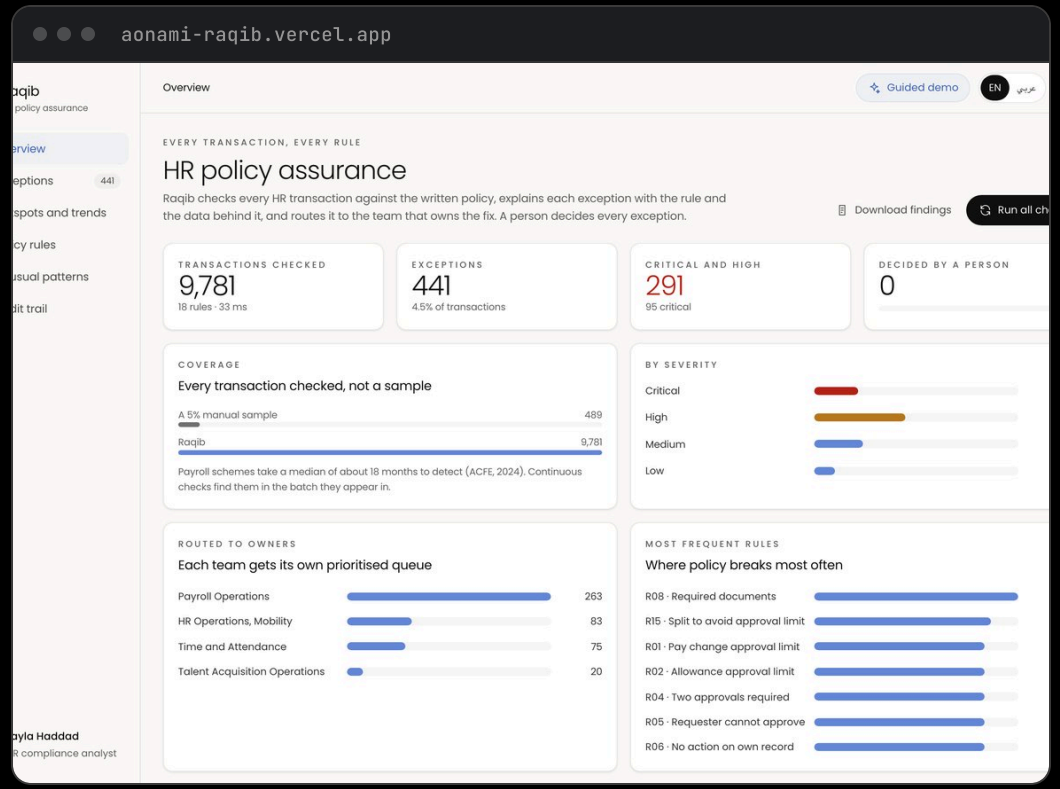
Raqib

Every HR transaction, checked against every rule. A person decides every exception.

LIVE PROTOTYPE

aonami-raqib.vercel.app

Press Guided demo. Runs on synthetic sample data.



Compliance by sampling finds problems late, if at all.

HR transactions are checked against policy by sampling. Approval limits, segregation of duties, required documents and timing windows are tested on a few hundred items, months after they happen.

TIME TO DETECT PAYROLL FRAUD

18 months

Median for payroll schemes, ACFE 2024. [1]

PAYROLL SCHEMES

~10%

of occupational fraud cases. [1]

CHECKED BY RAQIB

100%

of transactions, every batch, against every approved rule.

AT EMIRATES' SCALE

The Emirates Group workforce grew 8% to 130,919 people in 2025-26, with more than 4,000 UAE nationals. [5]

HOW THE CHALLENGE IS JUDGED

JUDGING CRITERION	WEIGHT
Detection accuracy	30%
Explanation and routing	25%
Responsible AI and security	20%
Deployability and reuse	15%
Team and pilot readiness	10%

Source: challenge brief [2]

A rules engine, not a guess.

Raqib replaces sampling with a deterministic engine: every transaction against every approved rule. The same batch always gives the same findings, each traceable to a rule and the exact data that broke it.

1 Ingest
Transactions and employees; columns matched by meaning.

2 Rules
Written policy as structured checks, each approved by a person.

3 Check
Every transaction × every rule, point-in-time data.

4 Explain
Plain English and Arabic, citing rule, fields and values.

5 Route
Owner by process, severity, suggested fix.

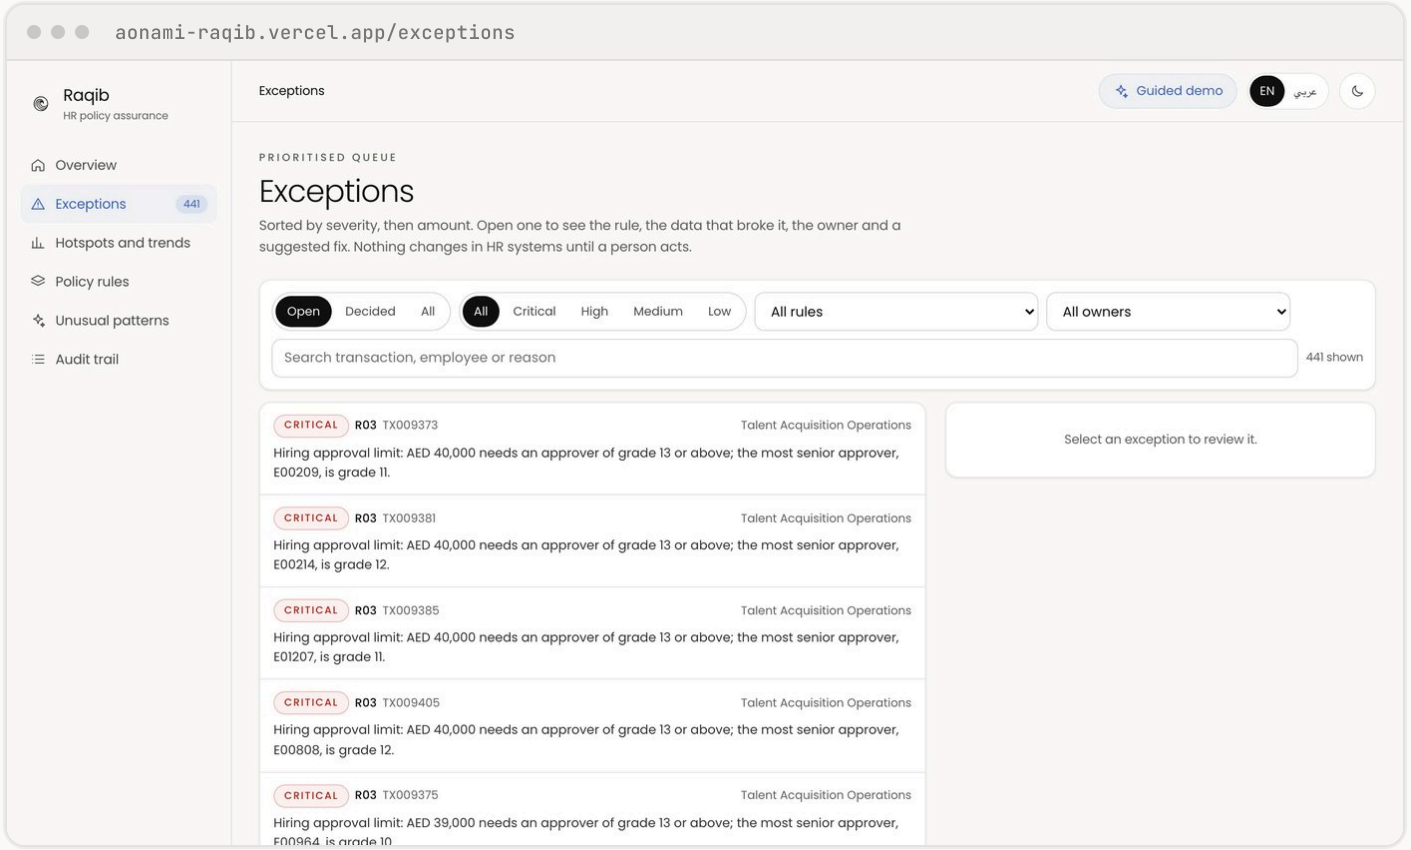
6 Decide
Accept, dismiss with reason or reassign; all audited.

- **No AI model decides** whether something is a breach.
- **Unusual patterns** are kept in a separate, unscored lane.
- **English and Arabic**, right to left, throughout.

AI WHERE IT HELPS

Detection is deterministic. Plain-English rules are matched against known rule patterns, no AI model, and a sentence that fits no pattern is refused, not guessed. An on-premise model could later suggest rule wording; it would never decide a breach.

One queue, critical first, each exception explained.



The rule, in one sentence

AED 40,000 needs an approver of grade 13 or above; the most senior approver is grade 11.

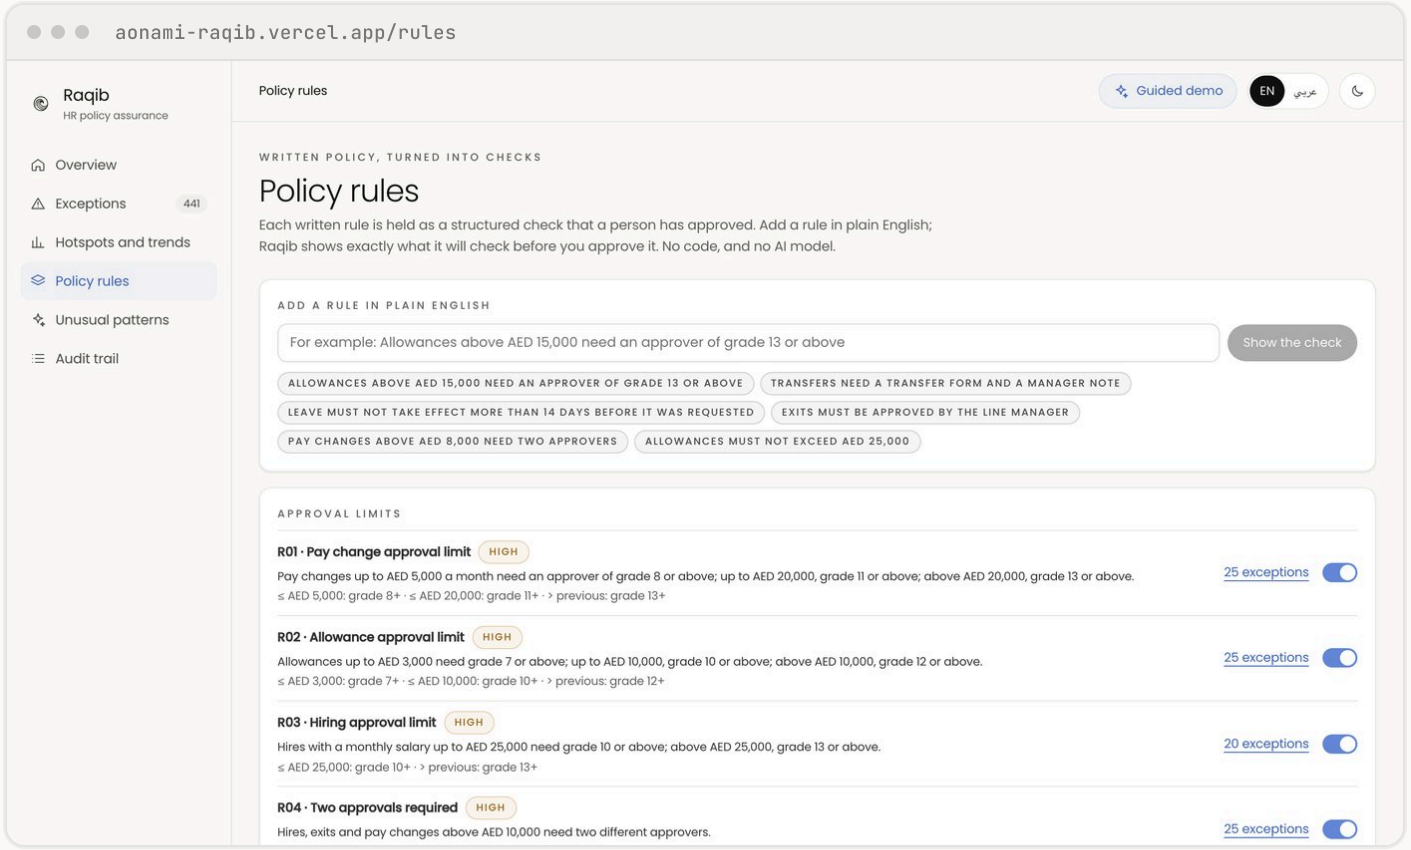
The evidence

Every field and value that broke the rule, so it can be checked in seconds.

A person decides

Accept and route, dismiss with a reason, or reassign. Raqib never changes a transaction.

Policy owners add rules in plain English. No code.



Plain-English rules

"Allowances above AED 15,000 need an approver of grade 13 or above" is matched to a known rule pattern and becomes a structured check. No AI model.

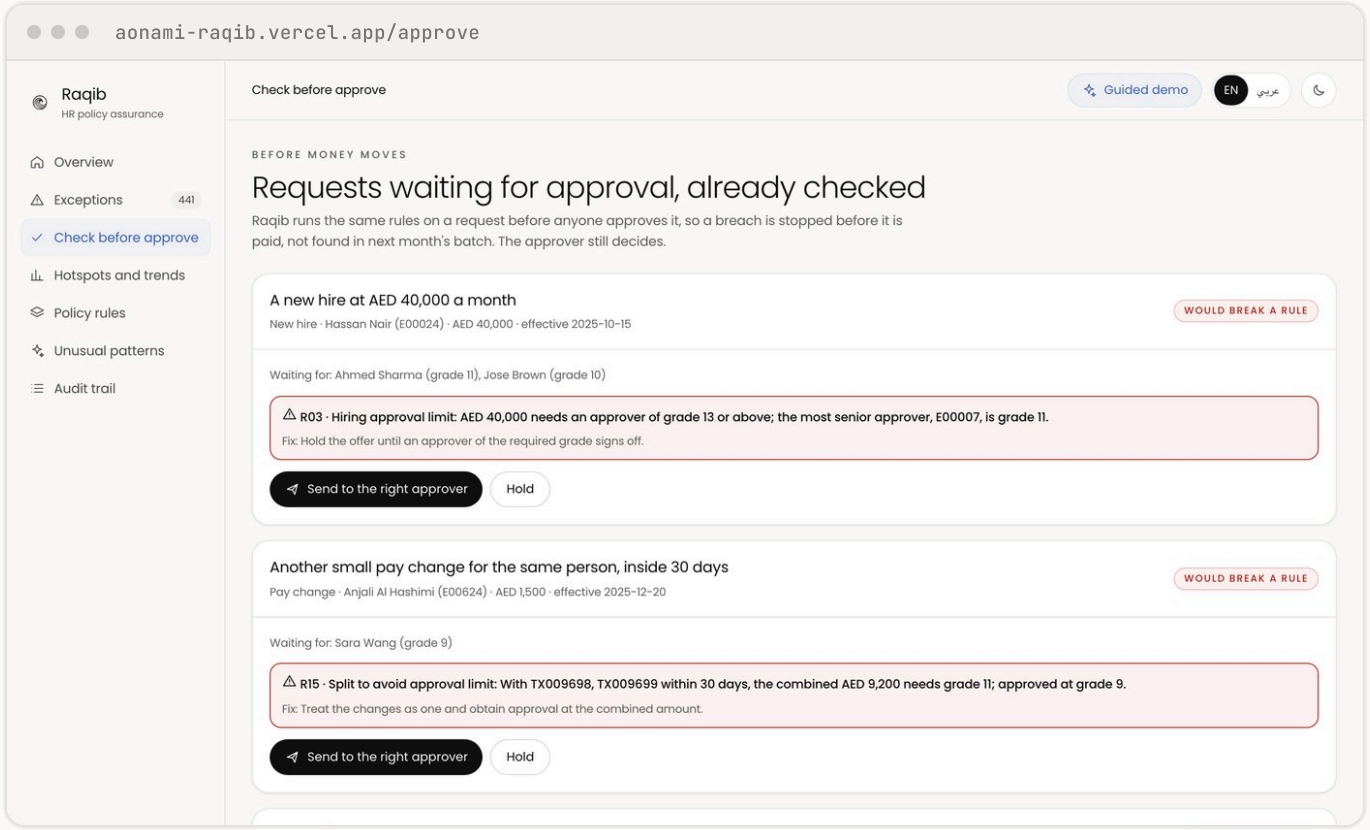
Approved before it runs

A named person approves each rule; the change is recorded.

Hotspots by team, never individuals

Where policy breaks most, by process, team and rule, month by month.

Check before approve: the same rules run before anyone signs off.



- A **AED 40,000 hire** waiting for grades 11 and 10 is stopped: it needs grade 13 (rule R03).
- **A split is caught at the source:** small pay changes for one person inside 30 days are added up before the next one is signed (R15).
- Same engine and rules, checked in milliseconds against the person's history. **The approver still decides.**

WHY IT MATTERS

Sampling finds a breach months later. A monthly batch finds it the same month. Checking before approval means it is never paid, across a workforce of 130,919 people. [5]

Built against a hidden answer key, the way it will be scored.

Before the official pack, we built our own: 2,000 employees and about 9,800 transactions over 12 months, 18 rules across the four named families, about 440 planted violations, plus 170 deliberate near misses that must not be flagged (exactly at a limit, one day inside a window).

LIMITS WE DISCLOSE

Results are on our synthetic benchmark. Accuracy on Emirates' rules depends on agreeing the exact meaning of each written rule with policy owners, which is the first step of the pilot. Arabic text needs native review.

WHAT WE EXPECT ON REAL DATA

Most first-run differences will come from how a written rule is read, for example which approver counts when there are two, not from detection. We settle each rule's meaning with its owner in weeks 1 to 2, then re-score.

BENCHMARK ON OUR OWN SYNTHETIC PACK	RESULT
Precision, three generated packs	99.8% to 100%
Recall, three generated packs	100%
Near misses wrongly flagged	0 of 170
Same data, different layout	Identical findings
Run time, 9,781 × 18 rules	well under 1 second

Our own synthetic packs; the official test file is the real test.

Every finding says what broke, why, who fixes it, and how.

R15 HIGH PAYROLL OPERATIONS

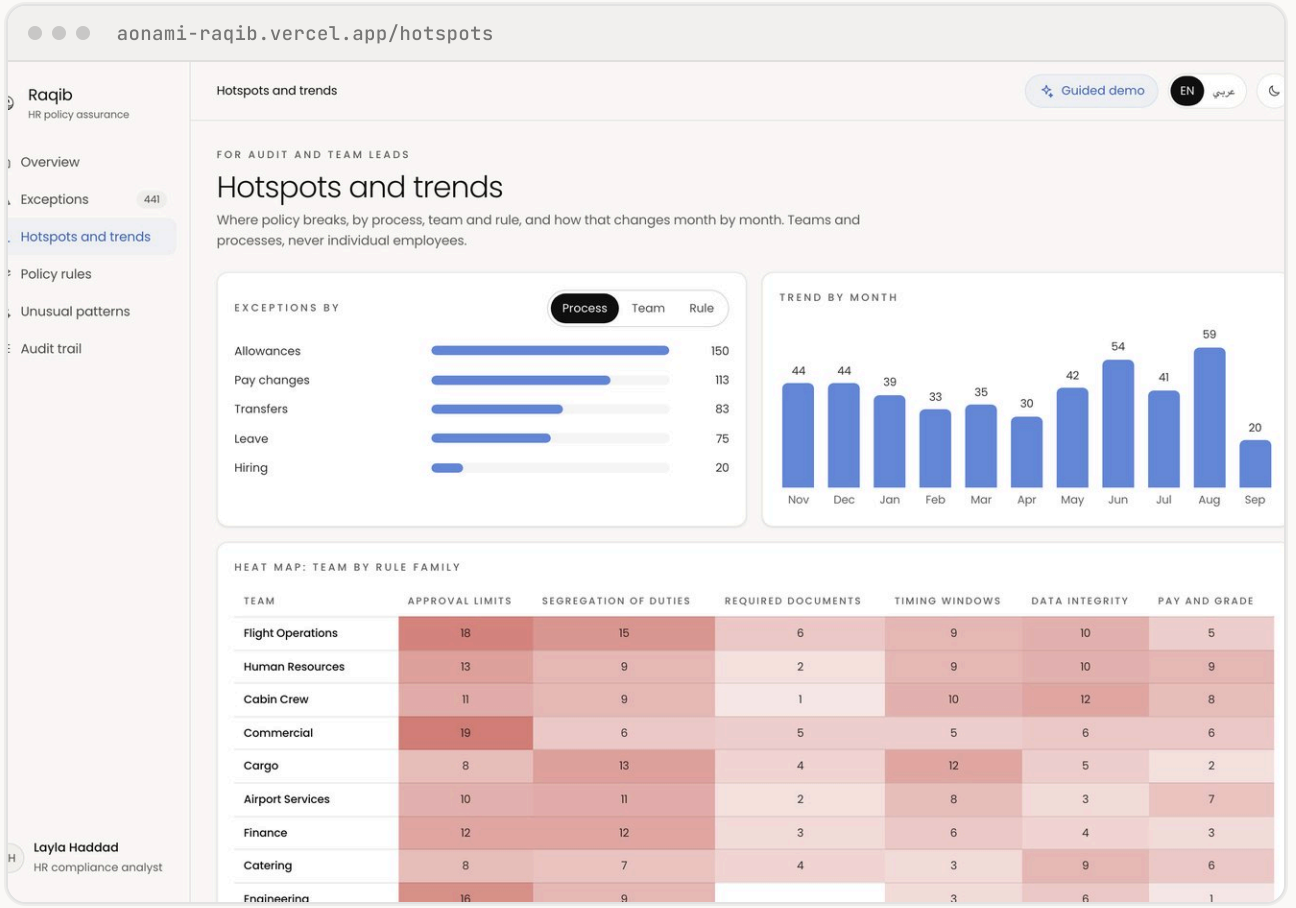
Split to avoid approval limit

With TX009158 within 30 days, the combined AED 9,100 needs grade 11; approved at grade 8.

`related=TX009158 (AED 5,000) · combined=AED 9,100 · required grade=11+ · approver grade=8`

Fix: treat the changes as one and obtain approval at the combined amount.

- Owners follow the process: payroll, hiring, mobility, time and attendance, offboarding.
- Severity starts from the rule and rises with materiality.



Raqib recommends. A person decides.

HUMAN DECISION ON EVERY EXCEPTION

Accept, dismiss with a reason, or reassign. Raqib never changes a transaction.

NO AI MODEL TOUCHES THE DATA

Deterministic checks; every flag traceable to a rule and the data behind it.

FULL AUDIT TRAIL

Every run, rule change and decision, with the person and the time.

NEUTRAL WORDING, TEAM-LEVEL VIEWS

“Does not match rule R07”, never accusatory. Hotspots by process and team, never by individual.

DATA STAYS IN EMIRATES’ ENVIRONMENT

Minimal fields; audit records hold IDs rather than personal details.

EMPLOYEES CAN RESPOND

In a pilot, anyone affected by a correction is told what changed and why before any recovery.

Designed for the UAE Personal Data Protection Law [3] and the principles of the UAE Charter for the Development and Use of AI [4].

Runs inside Emirates' environment. Another entity loads its own rules and nothing else changes.

1 HR system of record

Read-only exports or API pulls: transactions, employees, approval matrix.

2 Raqib

Static web app and command-line runner in a container. No outbound connection.

3 Single sign-on

Roles: compliance analyst, team lead, internal audit (read only).

4 Case or ticketing tool

Findings exported in the agreed audit finding format.

	HOW
Rules	Data, not code. Add or change a rule in plain English; approved by a person.
Other entities	Load their own policy and approval matrix; the engine is unchanged.
Schedule	The same engine runs in the browser and as a scheduled monthly batch.
Writes back	Never. Raqib only reads.

Six weeks, from policy to live monthly batch.

WEEKS	WHAT HAPPENS
1 to 2	Load the policy and approval matrix; agree each rule's meaning with policy owners
3 to 4	Shadow run on last quarter's transactions; compare with the sampling results
5 to 6	Live monthly batch with the compliance team; tune severity and routing
End	Coverage, exceptions found, time to decide; go or no-go

WHAT WE NEED FROM EMIRATES

- The written HR policy and the approval matrix
- A read-only export of last quarter's HR transactions
- A named policy owner per rule family and a payroll specialist
- A container host inside Emirates with single sign-on

WHAT WE MEASURE

- Coverage: share of transactions checked
- Exceptions found that sampling missed
- Time from transaction to decision
- Analyst agreement with each finding

THE DECISION AT THE END

Go or no-go, on numbers agreed with Emirates before the pilot starts.

Every judging criterion, and where it is answered.

CRITERION	WEIGHT	SLIDES	STRONGEST EVIDENCE
Detection accuracy	30%	6, 7	Breaches stopped before approval (check before approve); 99.8 to 100% precision, 100% recall; 0 of 170 near misses flagged
Explanation and routing	25%	4, 8	Rule, reason, evidence, owner, severity and fix on every finding
Responsible AI and security	20%	9	Deterministic, no AI model; a person decides; full audit trail
Deployability and reuse	15%	5, 10	Rules are data; plain-English rule authoring; container in Emirates' environment
Team and pilot readiness	10%	11, 13	6-week pilot with a shadow run against sampling

IN ONE LINE

Every transaction, every rule, explained and routed, and a person decides every exception.

TRY IT LIVE

aonami-raqib.vercel.app

Press Guided demo. Runs on our synthetic data.

Compliance tooling at scale, auditable data pipelines and measured AI.

JR Jayesh Kumar Rout

CO-FOUNDER, STRATEGY AND PRODUCT

IDfy: KYC and compliance, 60M+ verifications a month

OneCard: onboarding reporting engine for 6 banks

ON THIS BUILD

Owns the pilot plan and the relationship.

PK Pratik Kumar

CO-FOUNDER, PRODUCT AND GO-TO-MARKET

Stashfin: product manager, credit-report vertical

Aonami: designed the agent approval model in Bridge

ON THIS BUILD

Product flow and the demo; records the video.

BM Bhavik Munot

FOUNDING ENGINEER, PLATFORM AND DATA

Amazon: compliance-critical logs across 6,000+ AWS accounts

Electronic Arts: real-time pipelines up to 100K TPS

ON THIS BUILD

Data loading, matching and the command-line runner.

PJ Piyush Jha

FOUNDING ENGINEER, BACKEND AND INFRASTRUCTURE

Ideaflow, Menlo Park: founding engineer, prototype to production

Trell: APIs serving millions of daily users

ON THIS BUILD

Container deployment, single sign-on and connectors.

HA Harshil Agrawal

ML ENGINEER, EVALUATION AND MLOPS

Nielsen: multi-agent pipeline; 65% less manual analyst time

Monthly retraining and deployment for 50 clients, no manual steps

ON THIS BUILD

Benchmarks: the answer key, precision and recall.

Aonami builds AI agents and decision systems for regulated institutions, with a named person approving every action. Team of 9 across engineering, AI and delivery.

pratik@aonamitech.com

Every number in this deck, and where it comes from.

[1] ACFE, Occupational Fraud 2024: A Report to the Nations (payroll schemes in about 10% of cases; median 18 months to detect). https://www.acfe.com/-/media/files/acfe/pdfs/rttn/2024/2024-report-to-the-nations.pdf

[2] Challenge brief, Emirates Group on Ignyte Challenges (deadline 16 Oct 2026): judging criteria and weights. Provided to entrants; not public.

[3] UAE Government portal, Data protection laws: Federal Decree-Law No. 45 of 2021 on the Protection of Personal Data. https://u.ae/en/about-the-uae/digital-uae/data/data-protection-laws

[4] CMS Expert Guide, AI regulation in the UAE: UAE Charter for the Development and Use of AI (2024), principles incl. transparency, human oversight, accountability. https://cms.law/en/int/expert-guides/ai-regulation-scanner/united-arab-emirates

[5] Emirates Group, 2025-26 results (Dubai Media Office, 7 May 2026): workforce grew 8% to 130,919; UAE nationals above 4,000. https://www.mediaoffice.ae/en/news/2026/may/07-05/emirates-group-achieves-record-profit-in-2025-26

All people, records and figures on product screens are Aonami's own synthetic data. Emirates Group is named as the challenge host only; no endorsement implied.